PowerVR

by Imagination

PowerVR

Compute Development Recommendations

Filename	:	PowerVR Compute Development Recommendations
Version	:	PowerVR SDK REL_17.2@4910709a External Issue
Issue Date	:	30 Oct 2017
Author	:	Imagination Technologies Limited

Contents

List of Figures

List of Tables

1. Introduction

PowerVR Rogue is a family of Graphics Cores from Imagination Technologies designed to perform both compute and graphics tasks. Its programmable core architecture allows for highly efficient and high performance compute execution.

1.1. Document Overview

This document describes the recommended usage guidelines for achieving optimal performance when using Compute on Imagination's PowerVR Rogue Graphics Cores. Most of the guide describes in detail constructs and patterns that directly emerge from the PowerVR Rogue architecture, which should enable the developer to make the correct decisions irrespective of the preferred approach as far as APIs and programming languages are concerned. Additionally, several specific details are included for OpenCL, OpenGL and OpenGL ES Compute, and RenderScript. These details have been written against the following API versions:

- OpenGL ES 3.1
- OpenGL 4.3
- OpenCL 1.x and 2.0
- Renderscript

This document makes the assumption that the reader has a good working knowledge of at least one of these APIs. It is also assumed that the reader has worked through the examples of the relevant PowerVR Compute SDK or has referred to Google's documentation for RenderScript, available at:

- http://developer.android.com/guide/topics/RenderScript/index.html

After reading this document, readers should have a solid understanding of how Compute works on Rogue Graphics Cores, as well as a good understanding of how to develop efficient and well-optimised code for these devices.

Note: This document also provides an optimisation strategy quick reference sheet in Appendix C.

1.2. Glossary

Throughout this document, we will use the terminology identified in Table 1.

Table 1. Common terms

Term	Also referred to as…	Description
USC (Unified Shading Cluster)	Shading Cluster, Shading Unit, Execution Unit	A semi-autonomous part of the Graphics Core that can typically execute an entire workgroup. Other large parts like Texture Units can be shared among USCs.
Core	Processor, Graphics Core	An almost completely autonomous part of the Graphics Core. Typically, a collection of USCs and possibly supporting hardware such as texture units.

Term	Also referred to as…	Description
Task	Thread Group, Warp, Wavefront	The native grouping of threads that a USC executes. Consists of 32 threads on PowerVR Rogue cores.

OpenGL, OpenGL ES, OpenCL and RenderScript all use different terms to refer to, basically, the same concepts. In this document, we will be using what we consider to be the most unambiguous concepts, as listed in Table 2.

Table 2. API-specific terminology

Term	OpenGL and OpenGL ES	OpenCL	RenderScript
Kernel or shader	Compute shader	Kernel	Kernel
Thread	Thread, shader instance	Work-item	Kernel Invocation
Workgroup	Workgroup	Work-group	N/A
Shared memory	`shared` variables	`local` memory	N/A
Image	Texture, Image	Image	Image
Constant, constant memory	const/uniform variable, uniform block, uniform buffer	`constant` memory	constants
Private memory	Local variables, temporaries	Variables, `private memory`.	Local variables, Temporaries
Data set	Dispatch size, Data set	`Global work, ND Range`	Data set

2. A Modern System on Chip

As shown in Figure 1, a modern System on Chip (SoC) often integrates both a CPU and Graphics Core. The CPU is optimised for processing sequential, heavily branched data sets that require low memory latency, dedicating transistors to control and data caches. The Graphics Core is optimised for repetitive processing of large, unbranched data sets, such as in 3D rendering, dedicating transistors to arithmetic logic units rather than data caches and flow control.
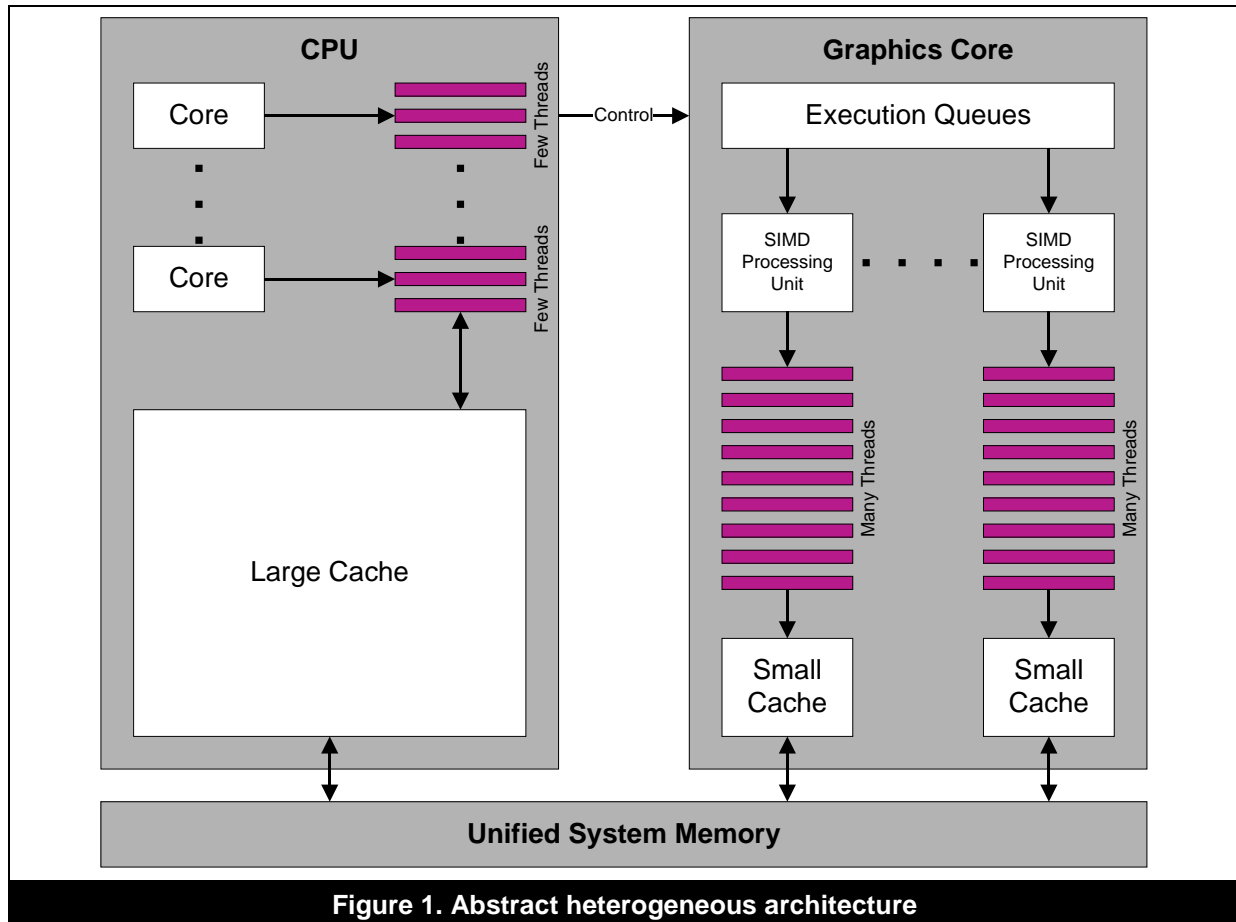
Figure 1. Abstract heterogeneous architecture

2.1.   Single Instruction, Multiple Data

Typical CPUs are optimised to execute large, heavily branched tasks on a few pieces of data at a time. A thread running on a CPU is often unique and is executed on its own, largely independent of all other threads. Any given processing element will process just a single thread. Typical numbers of threads for a specific program on a CPU is commonly one to eight, up to a few tens at any period of time.

Graphics Cores, on the other hand, work on the principle that the exact same piece of code will be executed on numerous multiple threads, often numbering into the millions to handle the large screen resolutions of today's devices, differing only in input and normally following the exact same steps, instruction for instruction. To do this efficiently, each processor executes the same instruction on multiple threads concurrently, in a form of Single Instruction, Multiple Data (SIMD) processing.

This should not be confused with vector processing, which is another form of SIMD. SIMD processors may either be scalar; operating on one element at a time, or vector; operating on multiple elements at a time.

2.1.1.   Parallelism

The main advantage of this type of architecture is that significant numbers of threads can be run in parallel for a correctly structured application and this type of parallelism is done with extremely high efficiency and effectiveness. SIMD architectures are usually capable of running many orders of magnitude more threads at once than a typical CPU.

Operating on large coherent data sets is something they perform exceptionally well at, which is why they are traditionally used for graphical and image processing. In actuality, algorithms that operate independently on a large coherent data set are well suited to this sort of processor.

2.1.2.   Divergence – Dynamic Branching

With each thread executing on one processor core, thread divergence has to be handled differently to how it works on a typical CPU. A CPU is generally made with the assumption of a sequential pipe – or workflow – that will be branching, looping, doubling back and otherwise following an extremely complex flow control in a very deep call stack. Today's CPUs are heavily optimised for this kind of work, with sophisticated branch prediction and other features aiming to optimise this kind of program.

On the other hand, a typical Graphics Core is optimised with the assumption that each "batch" of threads will execute the exact same code, instruction for instruction (lockstep), without any difference whatsoever. Additionally, it is expected, but not required that the input will only be "slightly" different (coherency).

Optimising the design for these assumptions leads to the already mentioned efficiency for these cases. The downside to these is that when branches are actually encountered (`if`, `while`, dynamic `for`), threads cannot each just execute a different path (it is physically impossible). Hence, the typical implementation is to have the processing core execute all paths of a branch in every thread, unless none of the threads take one of the paths, and just disable (mask out) all threads that have not taken the path. The actual mechanisms for this may be different on different architectures, but the net effect is the same.

In terms of computation, the cost of executing any branching code is the cost of executing all branches, unless none of threads hit the branch, plus the branching instructions themselves. It is therefore necessary to be careful of excessive branching when executing on a SIMD processor. There are, of course, exceptions to that: one of the most common and notable ones is conditionally loading or storing data from main memory, where it is always worthwhile to branch; usually, the bandwidth saved vastly outperforms the branch cost.

2.2.   Scalar or Vector Processing

Modern processing architectures can either operate on a single value per processing unit (Scalar) or multiple values per processing unit (Vector).

2.2.1.   Vector

Vector processing can be very efficient, as it can work on multiple values at the same time, rather than just one. For colour and vertex manipulation, this type of architecture is extremely efficient. Traditional rendering operations are, therefore, well suited to this type of architecture, as calculations often operate on 3 or 4 elements at once.

The main drawback of vector architectures is that if scalar values or vectors smaller than the processor expects are used, the additional processing element width is wasted. The most common vector width by far is 4, which means that a shader or kernel mainly operating on 3 component vectors will operate these instructions with 75% efficiency. Having a shader that does work on scalars in the worst case may take this number down to as low as 25%. This wastes energy, as parts of the processor are doing useless work. It is possible to optimise for this by vectorising code, but this introduces additional programmer burden.

2.2.2.   Scalar

Scalar processors tend to be more flexible in terms of what operations can be performed per hardware cycle, as there is no need to fill the additional processing width with data. Whilst vector architectures could potentially work on more values in the same silicon area, the actual number of useful results per clock will usually be higher in scalar architectures for non-vectorised code. For compute tasks it is thus a lot easier to work with scalar architectures, as there is no need to vectorise. Scalar architectures tend to be better suited to general purpose processing and more advanced rendering techniques. PowerVR Rogue cores are all scalar architectures.

3. Compute on the Rogue Architecture

Imagination's Rogue architecture is a programmable architecture capable of executing general purpose computations as well as rendering graphics.

3.1.   Architecture Overview

As shown in Figure 2, Rogue provides different hardware for transferring vertex, pixel and compute data between memory and the Graphics Core, and programmable processors for performing the actual operations. Unlike the Series5 architecture, Rogue has a dedicated path for compute tasks via the Compute Data Master (CDM), with the programmable arithmetic handled by Unified Shading Clusters (USCs).
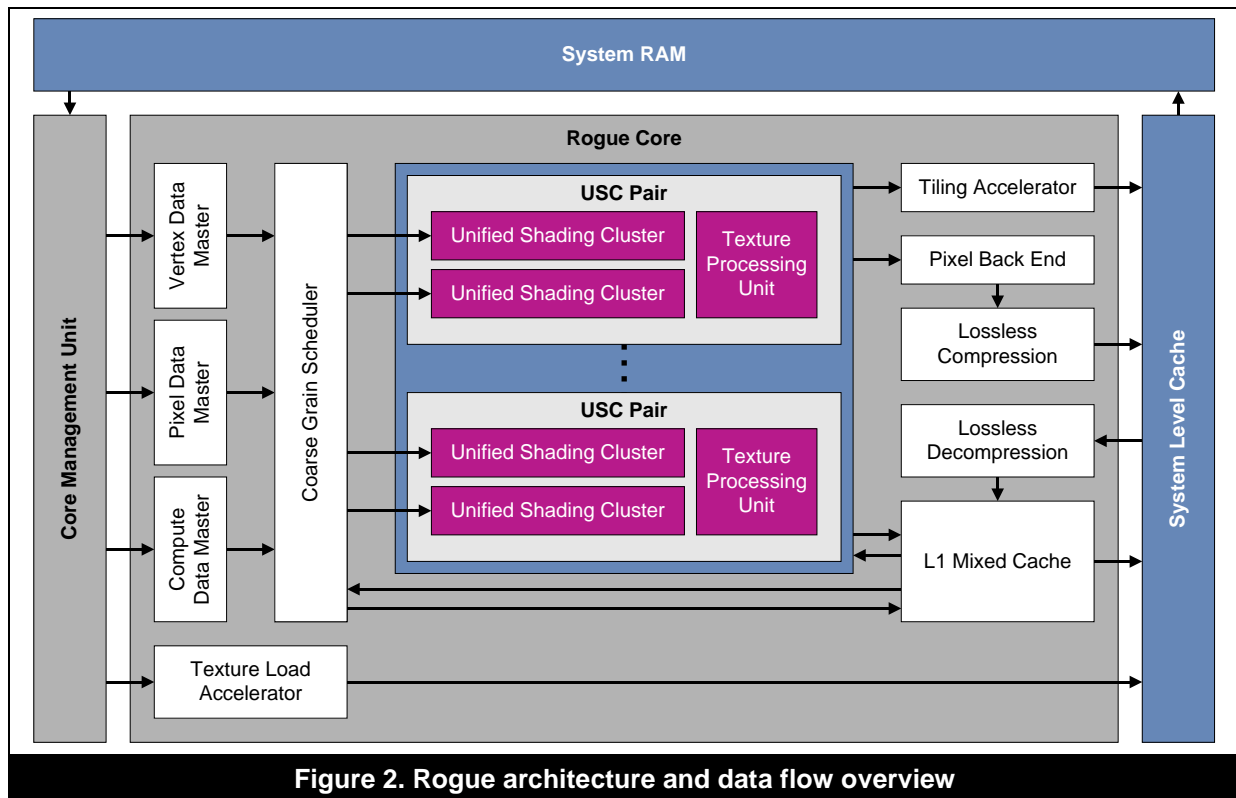
Figure 2. Rogue architecture and data flow overview

3.1.1.   Compute Data Master

Any compute jobs that have been dispatched by the host will need to be first translated into individual tasks for the Graphics Core to consume. The job of the Compute Data Master (CDM) is to take input from the host and generate these tasks.

3.1.2.   Coarse Grain Scheduler

Once tasks have been generated by the CDM, they have to be scheduled for execution. The Coarse Grain Scheduler (CGS) takes the generated tasks and on a broad scale schedules tasks across the available Unified Shading Clusters (USCs).

3.1.3.   USC Pair

Each USC pair contains two USCs and a Texture Processing Unit (TPU), to allow the most efficient balance between texture access requirements and compute density. The USC is the main processing element in the Rogue architecture.

3.1.4.   Texture Processing Unit

The TPU is a specialised piece of hardware used to accelerate access to images and textures from within kernel code. It handles image reads directly, with its own cache to ensure that image data transfers are as fast as possible. This hardware exists to better handle the typically huge size of images and the specialised access required by most image processing algorithms.

3.1.5.   L1 Mixed Cache

The L1 Mixed Cache is the main cache used by the Rogue architecture and all data transfers to and from memory go through here first. If this cache cannot serve a data fetch request, the request is passed to the System Level Cache.

3.1.6.   System Level Cache

The System Level Cache interacts directly with the System RAM and is the last chance for a data fetch to hit a cache. Data fetches that miss this cache are fetched from System RAM.

3.2.   Unified Shading Cluster

USCs are the programmable cores at the heart of the Rogue's computing power and are where the code in compute kernels is executed (Figure 3). The USC is a scalar SIMD processor, as described in Section 2. Each USC normally has 16 ALUs.
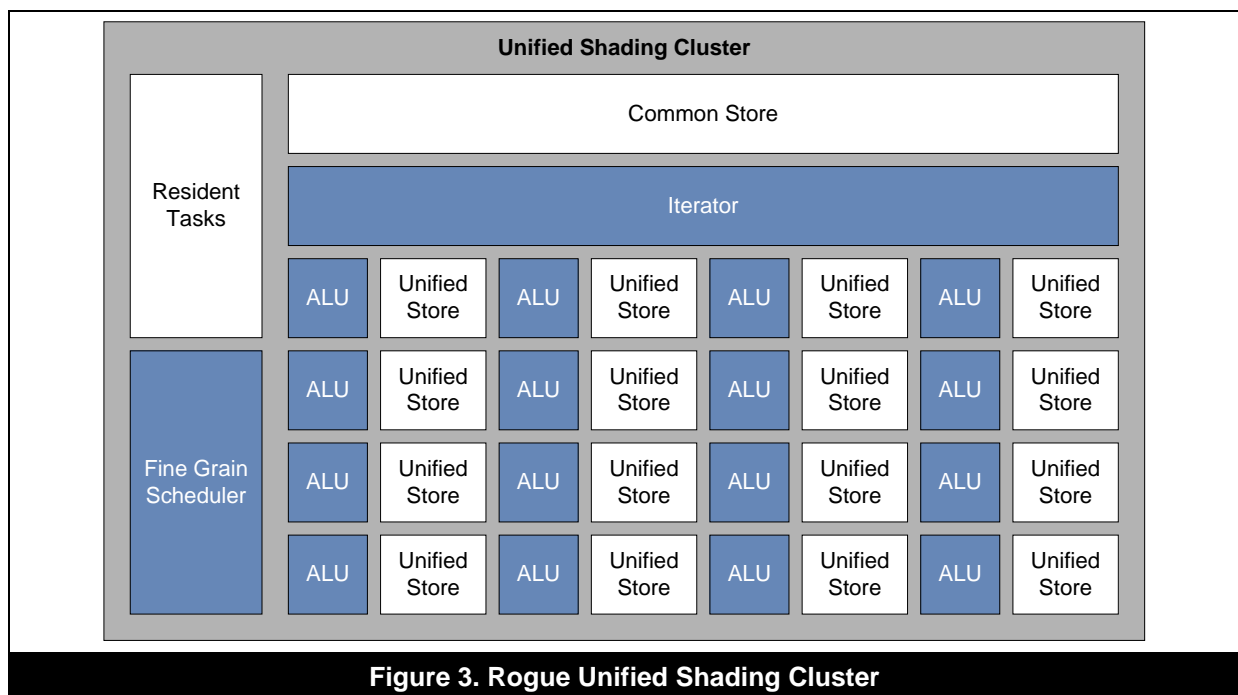
Figure 3. Rogue Unified Shading Cluster

3.2.1.   Execution

To execute the code within a compute kernel, the compiler has to translate the user's intentions into something the USC can understand. Equally, it is useful for the kernel writer to understand the mapping between hardware and API, to understand how to get the best performance.

Threads

Each work item of OpenCL, thread of OpenGL ES or kernel invocation for RenderScript, is handled as a thread of instructions in the hardware. These threads are executed together, in groups of (up to) 32 known as a task. In any given cycle, 16 threads (one per ALU Pipe) will execute the same instruction, meaning that half of a task runs in lockstep, with the other half staggered so that it executes one cycle behind.

Tasks

Tasks are the grouping used by Rogue hardware for SIMD execution of threads, with each task treated as a unit of work. Each task is executed by all ALU Pipes, one at a time, with half of the threads executing in lockstep in one cycle, and the other half executing in the next cycle.

In RenderScript the number of tasks used by a given execution is managed by the RenderScript compiler. If the global work size is a multiple of 32, the compiler will be able to separate threads to complete tasks, where every thread in a task is actively doing useful work. Any other number of items will leave threads vacant in some tasks, with tasks filling until all kernel instances are attached to a thread. The actual number of tasks and vacant threads per task will vary.

For OpenCL and OpenGL, while work-groups do not map directly to tasks, their size dictates task allocation and is arguably very important. Work-groups are split into tasks according to their size, as detailed in the following points:

- Work-group sizes that are a multiple of 32 allow each task generated to be fully utilised. Every thread in the task is actively doing useful work, with multiple tasks created until all work items are attached to a thread.

- Work-group sizes of 4, 8 and 16 typically execute with multiple work-groups per task – 8, 4 or 2 respectively. However, if there is barrier synchronisation in the kernel, and the work-group size was not specified at compile-time, tasks will execute one work-group with unused threads in the task left vacant causing a dramatic utilisation drop.

- Work-group sizes of 2 can at most be executed with 8 work-groups per task, utilising 16 of the 32 threads available, with 16 vacant threads.

- Other work-group sizes always lead to vacant threads, with each work-group only able to execute at most a single work-group.

Note: The maximum work-group size of an execution is 512. Higher sizes are unsupported.

Scheduling

The Fine Grained Scheduler (FGS) is responsible for scheduling and descheduling work on a per-task basis, from a pool of resident tasks. The size of this task pool will vary from device to device, but can typically be expected to be 48 or higher. The more tasks that are in this pool, the better any latency can be hidden.

One very important aspect of these tasks is that it is zero-cost for the USC: Tasks that need to wait for any reason are switched out and a ready to execute task is scheduled in its place immediately. This means that having large datasets scheduled at the same time generally allows greater efficiency without incurring scheduling overhead (unlike, for example, with CPU threads were scheduling cost is considerable).

Data Dependencies

When a data transfer dependency is encountered in a kernel, a task will need to wait for the data to be transferred before it can continue execution. The FGS will schedule out this task at no cost and schedule in another if one is available until the task is ready to resume.

Data dependencies are not just data reads and writes, but occur when the result of a read or write is required by another instruction. Any memory fences or barriers that occur within a kernel, or implicit fences caused by using the result of a read, will cause a dependency.

3.2.2. Memory

Memory space inside the USC comprises largely of two banks of registers: the Common Store which is shared across the entire USC and the Unified Stores, which are allocated per 4 ALU Pipelines.

Unified Stores

The Unified Stores are 4 small banks of registers in the USC, with each shared equally between 4 ALU Pipes. Any local/temporary variables used in a kernel or OpenCL `private` memory will be allocated into this space.

Each Unified Store consists of several 128-bit registers, with 1280 registers available in each. The actual number of registers available to a thread will vary in practice depending on a number of factors, but in general, using a maximum of 10 registers per thread allows the minimum recommended occupancy of 512 threads per USC.

These 128-bit registers are effectively accessed by a kernel as if they were a vector of 4 32-bit registers, so a kernel can use up to 40 32-bit values to hit the recommended utilisation. If a given kernel requires more than this, either the number of resident tasks goes down, or the number of threads executed per task does, so that each thread gets more register space. Imagination's compiler only reduces utilisation, as high occupancy is important to hide any latency in the threads.

For very excessive memory allocations above this threshold, utilisation can get to a point that it will not go any lower, and allocations will start spilling to main memory. This has high bandwidth and latency costs, though most applications will never hit this limit. The actual point that this happens is under software control and will vary from device to device.

Note: The compiler will try to reduce unified store allocations where possible, though it has to balance this with the instruction count use as well. For OpenCL the exact amount of register usage can be determined with a disassembling compiler which Imagination can provide under NDA.

Common Store

The Common Store is shared between all ALU Pipes in a given USC and is used for any data that is coherent between threads. Any shared memory (`shared` in OpenGL, `local` in OpenCL) is stored here, as are constant memory allocations (`const` or `uniform` in OpenGL, `constant` in OpenCL) that have known sizes at compile time. It is also used to store object handles, such as image, texture and sampler states, or pointers to main memory.

Data in the Common Store is allocated across 4 banks, each representing a column of data that is 128 bits wide, as shown in Figure 4. Data transfers can be served by multiple banks at once and an entire row of data can be read in a single cycle, or written in 4 cycles. Up to 16 threads in a task can fetch a 32-bit value each in one cycle if they fetch consecutive data from the same row (green).


Figure 4. Good (16 values/cycle) and worst (1 value/cycle) Common Store access

However, if more than one thread tries to access data from the same column (red) as another thread, it will cause a bank clash, as a bank can only serve one such request per cycle. Instead, it will serialise these fetches and it will take one cycle per-row to access the same amount of data, slowing down the pipeline. To avoid bank clashes, all data in the common store should be accessed across rows. Maximum throughput is, therefore, 512 bit/cycle: 4 128-bit registers (e.g., 4 32 bit values) from different banks each cycle.

For instance, imagine a kernel that uses a block of shared memory as a linear array of `float`. In the best case, if the stride between each kernel's access was exactly 32 bit (each thread fetches a subsequent `float` variable), all 16 threads would be able to fetch one value each in one cycle. The worst case scenario is that, if the stride between each kernel was 512 bit (kernels try to access every 16^{th} float in the array), the shader would need 16 cycles, one for each value, to complete the fetch. Formally, we could express the rules as:

- 4 banks of 128-bit registers.
- Each read can fetch an entire 128-bit register.
- Each 128-bit register can hold up to 4 32-bit values.
- Each read takes 1 cycle.
- No two registers from the same bank can be read per cycle (registers from the same bank cannot be read in the same cycle).

Example:

```
//OpenGL Compute
shared float common[SIZE];

int gid = get_global_id(0);

//Best scenario : 32-bit stride( 1 float)
//100% throughput (512 bit/cycle), no clashes. 1 32-bit value read/thread in 1 cycle.
float temp = common[gid];

//Worst scenario: 512 bit (16 float or 4 vec4) stride
//6.25% throughput(32 bit/cycle), 16-way clash. 1 32-bit value read/thread in 16 cycles
float temp = common[gid * 16];

//Normal scenario : 128-bit stride (4 float or 1 vec4), requesting 128 bits/thread
//100% throughput (512 bit/cycle), 4-way clash for 16 threads
//throughput. 4 32-bit values read/thread in 4 cycles
vec4 temp = vec4(common[gid*4],common[gid*4+1], common[gid*4+2],common  [gid*4+3]);
```

```
//OpenCL
local float common[SIZE];
int gid = gl InvocationId.x;

//Best scenario : 32-bit stride (1 float)
//100% throughput (512 bit/cycle), no clashes. 1 32-bit value read/thread in 1 cycle.
float temp = common[gid];

//Worst scenario: 512 bit stride (16 float or 4 float4)
//6.25% throughput(32 bit/cycle), 16-way clash. 1 32-bit value read/thread in 16 cycles
float temp = common[gid * 16];

//Normal scenario : 128-bit stride (4 floats or 1 float4), requesting 128 bits/thread
//100% throughput (512 bit/cycle), 4-way clash for 16 threads because we have exceeded max
//throughput. 4 32-bit values read/thread in 4 cycles
float4 temp = (float4 )(common[gid*4],common[gid*4+1], common[gid*4+2],common  [gid*4+3]);
```

The rule of thumb to avoid these scenarios is thankfully simple: At any one time threads should read consecutive values at the same time without any gaps. The most common ways one could accidentally enter that scenario would be having blocks of data per thread, instead of interleaved data.

A good way to iterate an array that contains multiple values for each thread is:

- **GOOD**: `float x = index * stride_per_index + thread_id`

Conversely, the following is very probable to cause multi-way bank clashes and waste bandwidth:

- **BAD**: `float x = thread_id * stride_per_thread + index`

System Memory Access

System memory is the Random Access Memory (RAM) on the device (outside of the Graphics Core), used by all system components. This is the type of memory that is referred to when discussing "memory bandwidth" in general. As many different subsystems need to access this memory, accesses have to be scheduled between them, increasing the latency of memory fetches.

OpenCL memory marked as `global`, OpenCL `constant` memory but with a size not known at compile time, Images, OpenGL Shader Storage Buffer Objects and RenderScript data in Allocators are stored in System Memory.

Always remember that the System Memory budget is limited, and shared between all resources. Many applications' performance will be limited by this resource making it a prime target for optimisation.

Rogue hardware can perform burst transfers on data transfers to/from system memory. This allows multiple bytes of data to be transferred in one go, without waiting for other devices. The more bytes that can be transferred in one burst, the lower the total latency of all memory accesses.

System Memory transactions are cached. A cache line in the L1 or System Level Cache is 128 bytes and maps to a 64-byte aligned segment in system memory. Transactions larger than 128 bytes are broken down into multiple cache-line requests that are issued independently. A cache line request is serviced at either the throughput of the L1 or L2 cache, in case of a cache hit, or otherwise at the throughput of system memory.

The hardware resolves System Memory accesses in row order for a given task's threads, so accessing data in row order will increase cache locality. If a task contains 16 threads that all perform an 8-byte memory transfer, and the transfers resolve to accesses that all fall within the same aligned 128-byte cache line, the access latency is the sum of the time taken to service a single cache miss plus the time to service a further 15 cache hits. This memory can be read and written in transactions of 1, 2 or 4 bytes, or multiples of 4 bytes and must be aligned to on-chip memory appropriately. 1 byte transactions are byte-aligned, 2 byte transactions must be 2-byte aligned and all other transactions must be 4-byte aligned. For best performance, data should generally be accessed from the kernel in this way. The compiler will split an unaligned transfer into multiple aligned transfers, which will be slower.

If a kernel performs all of its reads and writes at the same location in a kernel, it is possible for the compiler to combine multiple reads and writes into a single burst transaction. A burst transfers among the Unified Store and Global memory can transfer a maximum of 64 bytes at once. This includes any reads or writes within the shader or kernel between the system memory and temporaries (or `OpenCL` `private` memory). In order to perform these burst transfers, global memory must be aligned to 4 bytes, which normally includes all data types or vectors equal to or larger than 32 bits, such as `int`, `int2`, `int4`, `float`, `float2`, `float4`, `double`, `short2`, `char4`, but <u>not</u> `char`, `char2`, `bool`, `short`.

Burst transfers, specifically between System Memory and the Common Store, can transfer different amounts of data, depending on the size of the work-group. Specifically for OpenCL, if the following conditions are met:

- Transfer memory between `local` and `global` memory (any direction).
- A work-group's size is >=17 and specified at compile time.
- The OpenCL function `async_work_group_copy` is used.

Then, a different burst transfer to the common store can be executed, which provides a maximum of 4096 bytes across the workgroup with a single instruction. In the typical cases of preloading data into shared memory, this path should always be preferred. Otherwise, burst transfers to the Common Store (shared memory) can still be used efficiently but not as much as the above method.

4. Performance Guidelines

PowerVR Rogue Graphics Cores are capable of efficiently executing a range of parallel processing algorithms. To achieve the best performance, it is important to follow the guidelines throughout this section as closely as possible. These guidelines are based on four main strategies:

- Execute tasks on the most appropriate processor.
- Minimise bandwidth usage.
- Maximise utilisation and occupancy.
- Maximise instruction throughput.

For all of these strategies, it is always a good idea to profile applications to find out where the bottlenecks in processing are. This will work as a guide to where to target optimisation efforts and can be used to determine the results of any given optimisation. The PowerVR SDK provides PVRTune as a real-time hardware profiler which supports the visualisation of how the Graphics Core is being used.

4.1. Execute Tasks on the Most Appropriate Processor

If the hardware platform includes both a CPU and Graphics Core, an application should be structured so that it exposes both serial and parallel tasks. Serial tasks are most efficiently executed on a CPU, whereas parallel tasks are good candidates for executing on a Graphics Core.

To maximise system performance, tasks should be running on the CPU and Graphics Core at the same time, for example, by using the CPU to prepare the next batch of data while the Graphics Core processes the current batch. The Graphics Core is often a more power efficient processor if the algorithms used can be expressed as a highly parallel task. If the aim is to reduce power consumption, it becomes important to consider whether an algorithm can be expressed as a parallel task for execution on the Graphics Core, letting the CPU idle for this time.

4.1.1. Choose a Suitable Device

This is very different depending on the API. Firstly though, one should consider the work to be split between the CPU and Graphics Core; in other words between tasks suited for sequential or parallel execution. Tasks for the CPU can be implemented in any language supported by the platform, such as C++ or Java.

In OpenGL and OpenGL ES, tasks for the Graphics Core are normally dispatched as Compute Shaders. There are some cases where suitable alternatives exist and can be successfully implemented, such as using fragment shaders for image processing when displaying directly to screen or implementing time-evolving physical systems with vertex shaders and Transform Feedback. This document in general assumes Compute Shaders, as they are by far the most flexible of the three with the biggest area of application.

For OpenCL, the Installable Client Driver (ICD) enables multiple OpenCL drivers to coexist under the same system. If a hardware platform contains OpenCL drivers for other devices such as a CPU or DSP, an application can select between these devices at run-time.

The RenderScript runtime dynamically chooses whether a task can run on CPU or Graphics Core, depending on the workloads of each system and the suitability of the kernel for execution on each. So a developer cannot explicitly choose Graphics Core execution, but by using RenderScript and noting any device restrictions, it gives the runtime the choice to do so.

4.1.2.   Identifying and Creating Work for the Graphics Core

The USCs rely on hardware multithreading to maximise utilisation of their ALUs. As described in Section 3, the USC can execute instructions from a pool of resident threads, switching between them with zero-cost.

Keep the USCs Busy

The most common reason a thread is not ready to execute is that its next instruction is a pending memory access, which can span hundreds of cycles of latency. If other threads are schedulable (i.e., ready to execute) during this entire period, the memory latency is completely hidden, resulting in full utilisation of the USC (otherwise, completely eliminating the cost of the memory transfer).

If all threads are waiting for memory accesses, the USC will idle, reducing its efficiency (in other words, the application is bandwidth limited). This usually happens if the ratio of arithmetic to memory instructions in a kernel is low, or there are not enough threads available to hide the latency. Other system operations could also be reducing the bandwidth available.

If an application is proving to be bandwidth limited, it is worth trying to reduce the number of memory fetches in each kernel, or favouring run-time calculations over lookup tables. More work items can also be provided to help mask this. Tools such as PVRTune can be used to help detect when this happens.

Avoid Overly Short Kernels with Small Data Sets

For a given kernel execution, the first few threads often run with reduced efficiency, due to the time taken to allocate work to the USCs. The last few executions also often run with reduced efficiency as the pool of resident tasks empties. The more work that is done per execution, the lower the percentage of this cost:

- Use somewhat longer kernels, which can mask out the latency of the scheduling/descheduling.
- Work on larger data sets, which create more threads per execution that can all be scheduled together. Typically, this should be in the order of several thousands to millions of data points.

If an algorithm has a small kernel length or works on a small data set, it may be difficult to achieve high USC efficiency. In this case, the algorithm may be more suited for execution on a CPU.

4.1.3.   Sharing the Graphics Core between Compute and Graphics Tasks

A long-running kernel can starve system components that rely on 3D graphics processing, such as a user interface. In this situation, a long-running kernel can be split into multiple kernels executed one after another. Note that this typically requires the first kernel to write some data to global memory and the second to read this data back, which will introduce additional bandwidth usage.

4.2.   Minimize Bandwidth Usage

If an application does not need to access memory or move data around, then it should not. Several changes can allow an application to avoid this depending on the algorithm used, including doing calculations at run-time instead of using a lookup, or reducing the number of memory fetches for an image filter. Other optimisations can be performed for Rogue Graphics Cores as laid out next.

4.2.1.   Avoid Redundant Copies

One of the first steps in reducing memory throughput for an application is to ensure that hardware components using the same memory all access the same data, without any intermediate copying required. Examples of hardware components include the CPU, Graphics Core, camera interfaces and video decoders.

Create Allocations with Correct USAGE Flags

Applicable to: ☑ RenderScript

Rogue devices are capable of allocating memory and then mapping the allocation to a CPU pointer, but cannot map arbitrary host-allocated memory into the Graphics Core. In RenderScript, device

memory models are treated semi-opaquely, with no explicit indication to the developer whether memory is allocated for the CPU or Graphics Core. Memory spaces are defined for Allocations via "USAGE" flags however, with each usage being treated as a potentially separate backing store.

Allocations should be explicitly synchronised across usages each time they are used in a different context, to ensure that everything has the same view of memory. Synchronisations can cause copies to occur when memory is allocated in different memory spaces and so should only be used when necessary. If an algorithm can be designed to avoid synchronisations at any point, it should do so.

Create Shared Memory Objects with CL_ALLOC_HOST_PTR

Applicable to: ☑ OpenCL

In OpenCL, host and device memory models are specified independently from each other, with interaction performed by either explicitly copying data or by sharing regions of a memory object. Rogue devices are capable of allocating memory and then mapping the allocation to a CPU pointer, but cannot map arbitrary host-allocated memory into the Graphics Core. Therefore, memory objects that are to be shared should be created using the flag `CL_ALLOC_HOST_PTR` and accessed using mapping functions `clEnqueueMapBuffer` or `clEnqueueMapImage`, and `clEnqueueUnmapMemObject`.

If an object is created using the flag `CL_USE_HOST_PTR`, prior to enqueueing a kernel, the driver will create a copy of it that can be mapped into the Graphics Core address space. This copy occurs every time data is transferred between the host and device and will increase system bandwidth demand.

Create Buffers with the Correct Usage Hints

Applicable to: ☑ OpenGL ☑ OpenGL ES

In OpenGL and OpenGL ES, buffer allocation is handled by the driver semi-opaquely with usage hints during buffer allocation. The same recommendations apply as with OpenGL and OpenGL ES and it is, therefore, necessary to use the hints that are correct for one's application. Moreover, in general, mapping solutions (`glMapBufferRange`) are preferred to explicit uploads (`glBufferSubData`) when the data are going to be changed "frequently".

Use Zero-Copy Paths if Possible

Applicable to: ☑ OpenCL ☑ OpenGL ☑ OpenGL ES ☑ Interop

When processing input data from an external source such as a camera module, a zero-copy path can be enabled between the camera data and the API image used by the Graphics Core, by using `EGLImages`. This typically means directly rendering an image - see the example in Figure 5 for communicating an image from the Camera module.

Figure 5. Zero copy from a camera to OpenCL and from OpenCL to OpenGL ES
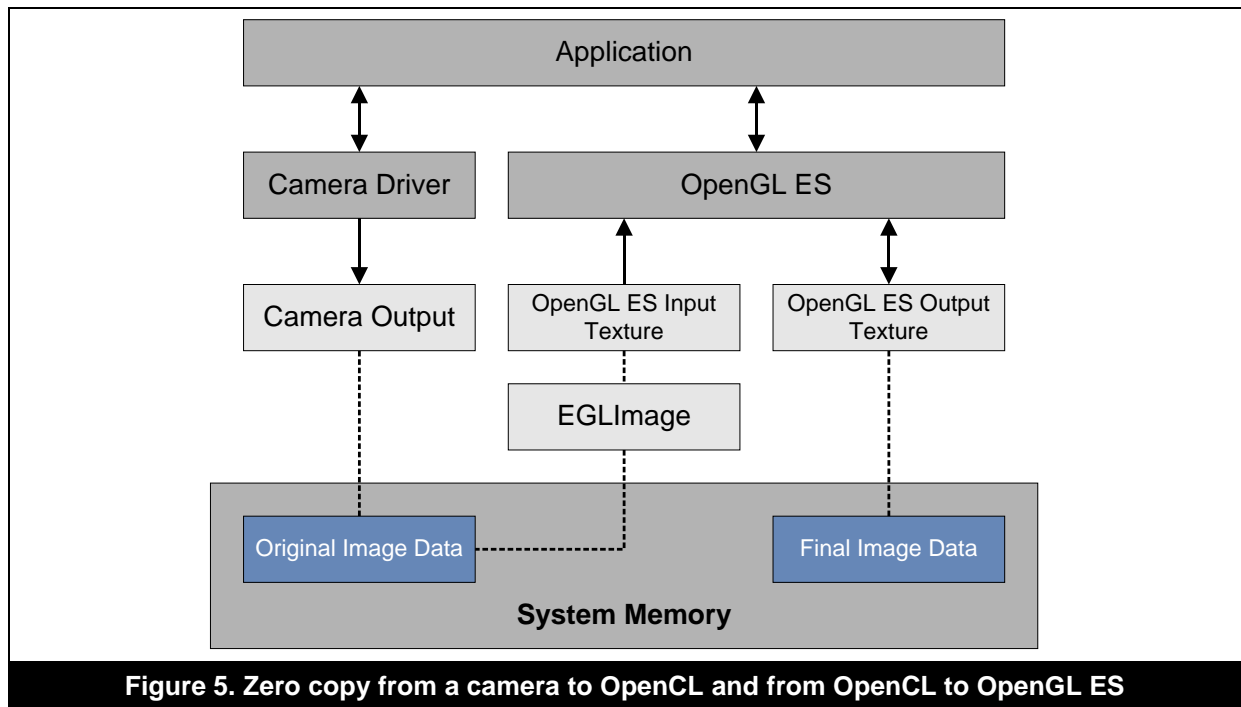
OpenGL and OpenGL ES

If the output image is to be used outside the graphics pipeline, it can also be shared with an `EglImage`. If the data is not required outside the graphics pipeline, it may be worth it to consider directly using a fragment shader for rendering directly to the framebuffer instead of using Image Load/Store.

OpenCL

A zero-copy path can also be enabled between the OpenCL output image and the OpenGL ES input texture used for rendering by using EGL Images (Figure 6). An example of how to do this (TextureStreaming) is included in the PowerVR SDK, alongside a whitepaper explaining how it works.

Figure 6. Zero copy from a camera to OpenCL and from OpenCL to OpenGL ES
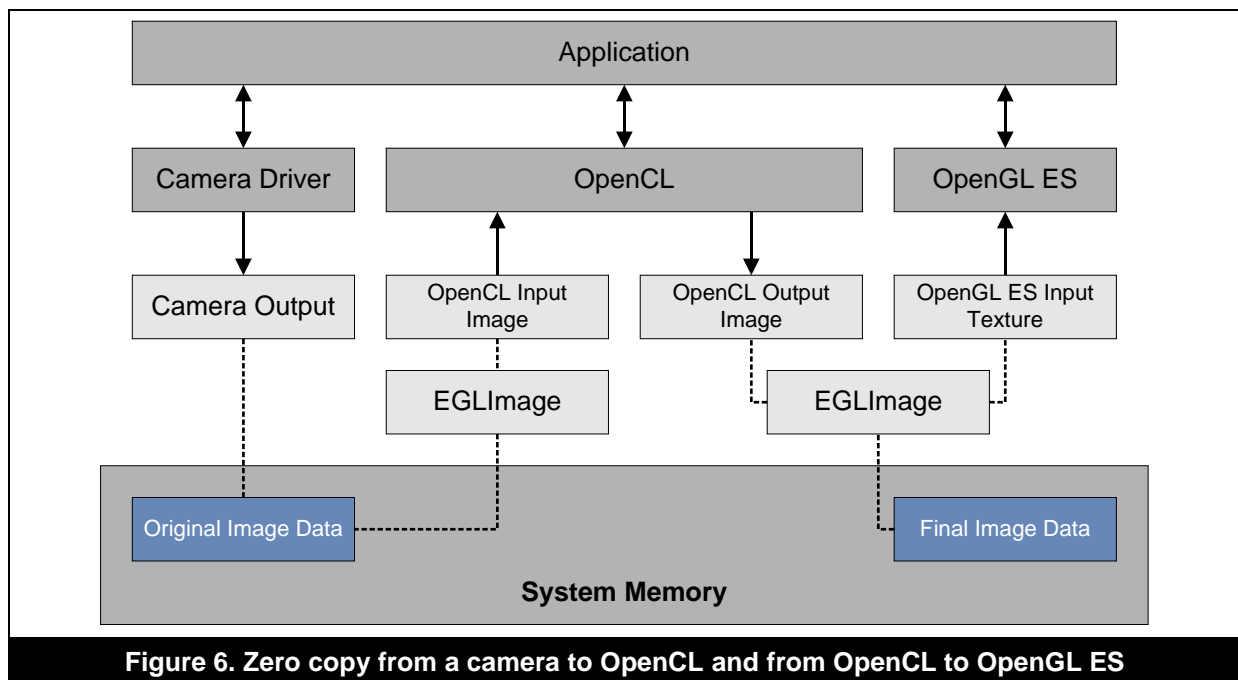
4.2.2.   Group Memory Accesses Together

The compiler uses several heuristics and can identify memory access patterns in a kernel that can be combined into burst transfers for read or write operations. To allow this to happen, memory accesses should be grouped together as closely as possible in order to be as easy to identify as possible. Generally, putting reads at the start of a kernel and writes at the end allows for the best efficiency. Accesses to larger data types such as vectors also compile into single transfers wherever possible, which implies that loading a single `float4` is preferred over 4 separate float values.

4.2.3.   Use Shared/Local Memory Sensibly

Applicable to: ☑ OpenCL ☑ OpenGL ☑ OpenGL ES

Unlike Series5, Rogue devices have a good amount of shared memory available to make good use of workgroups. Kernels should generally prefetch data into shared memory if it is going to be accessed from multiple kernels. This is effectively a software-controlled cache, augmenting the hardware controlled caches on the device.

Typically, the way to make good use of this is to evaluate how many accesses are going to be needed across work-items, then cache these values into shared memory at the start of a kernel. This work should ideally be done with explicit commands where available (OpenCL `async_work_group_copy`), or otherwise be split between available work items as evenly as possible, with no two threads caching the same values. Barriers should be used to synchronise all work-items as late as possible, but before using the values. Then, each kernel should proceed as usual, using these cached values.

4.2.4.   Access Memory in Row-Major Order

To make best use of the hardware's memory architecture, all memory should be accessed in row-major order across threads. This applies to all memory banks in the architecture, though for subtly different reasons as described in Section 3.2.2. Generally, memory can be considered as linear, even when allocated as an N-Dimensional array. This means that two contiguous values in the last dimension are physically next to each other in memory, and can often be accessed together in a single transfer. This is much more important when working in mobile devices with limited bandwidth on large data sets.

Note about OpenGL/Images: This applies to most kinds of memory and buffers, but specifically does not apply to texture or image objects. These are optimised for cache efficiency based on N-

Dimensional access and dedicated hardware handles this (Section 3.1.4). In these cases, it is important to access nearby pixels to ensure cache coherency. For example, a typical 2 dimensional texture should generally be accessed in square/neighbouring pixels starting from (0, 0) and increasing.

4.3. Maximise Utilisation and Occupancy

Utilisation is the efficiency of execution for each task, with regards to the number of ALU Pipes being used at any given time. If a given task does not have a full 32 threads performing useful work, then available processing power is being wasted and utilisation is low (one should always strive to keep hardware busy).

Occupancy refers to the number of tasks that are queued, ready for execution in a USC. The FGS can swap tasks seamlessly if a task is already resident, which hides any latency that would otherwise be caused by stalls in a task that is currently running. Section 3.2.1 gives more details on how the hardware manages tasks and threads.

4.3.1. Work on Large Data Sets

The most important part of ensuring high utilisation and occupancy is to ensure there are enough data points to make best use of the hardware. Data sets larger than about 512 items per USC on-device typically provide enough work to maintain high utilisation and occupancy, with larger numbers of items increasing efficiency further.

Aim for a Data Set Size that is a Multiple of 32

If a data set has a size that is a multiple of 32, there is the best opportunity for full utilisation. While for sufficiently large datasets the difference in utilisation should be negligible, care needs to be taken when doing this to ensure that either filler data do not affect the final result, or with code in the kernel. For OpenCL, a kernel's global work size can be set to a padded size, whilst the actual data set is kept at its original size. In this case, the kernel has to be careful to avoid out of bounds memory accesses.

4.3.2. Choose an Appropriate Work-Group Size

The number of ALU Pipes active in any given cycle depends on the number of threads assigned to each task. A task consists of 32 threads and a few things affect whether these are all used fully or not. In OpenCL and OpenGL ES Compute a developer has fairly precise control over this by choosing a work-group size, so it should be chosen carefully.

Aim for a Work-Group Size of 32

Ideally, a work group size of 32 is perfect for execution on Rogue hardware, as this exactly matches the size of a task. This ideal situation is not always possible however, depending on the size of the total data set actually being worked on and the local problem space. Suggested work-group sizes are listed below, in order of decreasing efficiency:

- 32.
- Multiples of 32, up to and including 512 (practically ideal).
- 16, 8 or 4, specified at compile time (only recommended if no alternative exists).
- Any other size (to be avoided – will have a number of threads per workgroup idle).
- 2 (to be avoided – half of the threads will be idle).

For a discussion on exactly how the compiler handles work-group sizes, see Section 3.2.1.

Tell the Compiler the Work-Group Size

Applications should always tell the compiler how large a work-group a kernel expects to work on.

For OpenGL and OpenGL Es this means using

`layout(local_size_x=X,local_size_y=Y,local_size_z=Z.`

For OpenCL this means using `__attribute__((reqd_work_group_size(X, Y, Z)))`.

By knowing the size, the compiler can better allocate on-chip resources, such as shared memory, and perform targeted optimisations at compile time. Deferring the choice until the kernel is enqueued or letting the runtime choose misses out on these opportunities.

In the case where a kernel is intended for use with multiple different work-group sizes, compiling the kernel multiple times with different work-group sizes is usually a much better option than deferring the choice to runtime. The performance gains can be quite significant in some cases and is usually worth the additional application complexity.

For work-group sizes of 16, 8 or 4 where barrier synchronisation is used, full utilisation can only be achieved if this attribute is specified.

4.3.3. Reduce Unified Store Usage (Private Memory)

The USC's unified stores are shared between all work-items that are resident on a USC at any given time and are used to store any temporary variables, or OpenCL `private` memory allocations, and in the kernel.

The compiler prioritises high occupancy over high utilisation and will reduce the utilisation of each task based on private memory only. Having at least 512 threads resident per USC provides the best opportunity to hide any data fetch latency, whilst maximising utilisation, which equates to a maximum of 40 scalar values used by the kernel.

If a kernel uses less private memory than this, more tasks can be made resident on most cores. The compiler will attempt to optimise unified store private memory use where possible, reducing the final amount. Under NDA, Imagination can provide a compiler which gives exact values for private register usage.

4.3.4. Reduce Common Store Usage (Shared/Local Memory)

`Shared` (OpenCL `local`) memory is extremely useful as it can act as a software-controlled caching mechanism, thus reducing bandwidth requirements in an application, as discussed in Section 4.2.2. A lot of optimisations usually start by caching into this store.

However, there are limits to how much can be used before affecting occupancy. As with Unified memory, resident workgroups share local memory on a USC. The more work-groups that can fit within the maximum available local memory, the greater the opportunity the core has to hide latencies in kernel execution.

4.3.5. Be Aware of Padding

To improve data access speed, all data types less than 32-bits are padded to 32-bits in both the unified and common stores. This means that a `char` or `short` actually takes up 32-bits of register space.

Note for OpenCL: When querying the amount of `local` memory available to a device, the number returned is how many 32-bit values can be stored, which applies equally to `float, int` or `char` values.

For some kernels this padding means that register use may start going beyond best performance limits, as discussed previously, when it would not if they were not padded. In these cases it may be useful to manually pack or unpack values from one or more 32-bit types. This allows values to be stored for their lifetime as smaller values and only being expanded when needed. There is a computation cost associated with this, so care should be taken.

4.4. Maximise Instruction Throughput

To get the best performance, the hardware needs to do as many calculations as it can in as short a time as possible. Ensuring the right algorithm is used is the most important thing to do here, as detailed in Section 4.1.2. Once the correct algorithm is chosen, a number of optimisations can further improve performance, as detailed below.

4.4.1. Trade Precision for Speed

Unless specifically needed for a particular scenario, precision should liberally be traded for speed, to minimise the use of arithmetic instructions with low throughput, as long as care is taken to ensure there is enough precision to produce acceptable results.

For OpenGL and OpenGL ES, consider using `mediump` floats whenever practical – these may have optimised paths by requiring less precision, resulting in some reasonable gains in some common instructions.

For OpenCL kernels, <u>always</u> use the `-cl-fast-relaxed-math` build option, as this enables aggressive compiler optimisations for floating-point arithmetic which the standard otherwise disallows. This flag will often result in good performance gains at the cost of very little precision. `native_*` math built-ins should also be used, as they are much faster than their un-prefixed variants for the cost of some arithmetic precision.

For RenderScript, enable the pragma operation `#pragma rs_fp_relaxed` to enable some optimisations, or `#pragma rs_fp_imprecise` to allow the compiler to perform more aggressive optimisations.

4.4.2. Tweak Work for the Graphics Core

The PowerVR Graphics Core has an impressive arithmetic throughput and fully supports integer and bitwise. Nevertheless, it is still beneficial to avoid instructions that are excessively long or complicated. It is important to remember that all Graphics Cores are optimised for floating point calculations and may achieve better throughputs with them.

The integer path is reasonably fast, especially in combinations of arithmetic (except division!), bit-shifts and when tests are needed, and several of those may be possible to process per cycle.

That said it will normally be faster still to work in floating point. Ideally, the numbers should be float in the first place. For kernels with a lot of arithmetic on comparatively few values per kernel, it may be beneficial to turn even integer values into float and back to integers after calculating. If in doubt, use a profiling compiler as these kinds of gains will be reflected in the cycle count of the kernel.

Integer division should always be avoided if at all possible. Even the worst case scenario of casting to float, dividing and casting back to an integer will be a lot faster than a true integer division. The caveat of that is of course the reduced range of input; not all integers can be represented by float, hence the results will be accurate for a specific range. The integer numbers that are exactly representable by a 32-bit float are up to roughly 16.7M, but the reciprocal involved may cost extra accuracy, so it is advised to be wary with numbers even in the low millions.

4.4.3. Avoid Excessive Flow Control

Any instruction that changes what code is executed based on some condition is considered to be flow control. Any flow control statements (`if`, `?`, `switch`, `do`, `while`, `for`) in a kernel may cause dynamic branching as discussed in Section 2.1.2.

PowerVR Rogue is very flexible and efficient with typical branching, but a branch is always a branch and inherently inefficient in a parallel environment. Branching should be avoided where practical, either with compile-time conditionals using the pre-processor or by replacing them with an arithmetic statement instead. Note that when conditionals are used to decide whether to access memory, the bandwidth savings will normally far outweigh the cost of the branch.

Also, note than when very few instructions are present in any path of a branch statement, predicates will often be used instead, which results in less instructions used overall. Whether this happens or not can be seen when using a disassembling compiler (available under NDA).

Alternatively to using actual branches, it is often possible to use built-in functions such as `clamp`, `min`, `max` or even casting a `bool` to a `float` or `int` value and performing arithmetic. Replacing branches with arithmetic can be tricky, but in some cases can result in decent performance boosts when done well, as these do not cause any kind of divergence among threads.

4.4.4.   Avoid Barriers Directly after Local or Constant Memory Access

Applicable to: ☑ OpenCL ☑ OpenGL ☑ OpenGL ES

The hardware may have a latency of up to two cycles for calculating an array index and using it to fetch shared memory from the Common Store (see Section 3.2.2). The compiler is generally able to hide this latency by rearranging instructions to carry out arithmetic during this time. If a barrier is present immediately after an access, as is often the case, then the compiler typically cannot do this. To allow the compiler to hide the latency, as much arithmetic that does not depend on the result of the access as possible, should be performed before the barrier is inserted.

4.4.5.   Use Built-ins for Type Conversions

Applicable to: ☑ OpenCL ☑ OpenGL ☑ OpenGL ES

While generally not very expensive and sometimes can be completely hidden, the compiler may sometimes require some additional instructions to convert between different data types. When suitable, explicit pack/unpack or conversion functions should generally be preferred to manual operations.

OpenCL

OpenCL actually provides a very clear and complete suite of type conversion functions. When they are present, one should use the `convert_type()` functions. When conversions are required, follow the conventions:

- For conversions from `int` to `float`: round to zero (`rtz`) is fastest.
- For conversions from `uint` to `float`: round to zero (`rtz`) and round to negative infinity (`rtn`) are fastest, round to positive infinity (`rtp`) is relatively fast and round to nearest even (`rte`) is slow.

Also useful are the `as_type()` functions, which are analogous to reinterpret casts in C++. These allow a kernel write to access two types that are conceptually the same size (e.g., `int` and `char4`) to be interpreted interchangeably. Types that are not actually the same size on-chip have a higher cost than those that do, as additional instructions are required to extract the correct bits. For example, a `char4` actually uses 128-bits of register space, whereas `as_type()` treats it as 32-bits. Packing instructions need to be used to compensate, as opposed to, for example, an `int4` treated as `float4` would typically be a no-op.

OpenGL and OpenGL ES

For OpenGL and OpenGL ES, the relevant conversions are a suite of `packXXXXXX` or `unpackXXXXXX`, and generally allow one to pack and unpack floating point values relevant to the graphics pipeline. Consult the specific API version for available instructions.

5. Contact Details

For further support, visit our forum:
http://forum.imgtec.com

Or file a ticket in our support system:
https://pvrsupport.imgtec.com

To learn more about our PowerVR Graphics SDK and Insider programme, please visit:
http://www.powervrinsider.com

For general enquiries, please visit our website:
http://imgtec.com/corporate/contactus.asp

Appendix A.   Measuring Peak Performance

Throughput, in Operations-Per-Second, is one of the most important factors for performance. The theoretical maximum throughput for an arithmetic operation can be expressed as the following equation:

$$T_{max} = \frac{\text{Clock speed} \times \text{Number of USCs} \times \text{ALU Pipes per USC} \times \text{Operations per ALU Pipe}}{\text{Cycles per Operation}}$$

Each USC has 16 ALU Pipes that run in parallel (see Figure 3). Rogue USCs are optimised to maximise the number of scalar floating-point operations that can be processed on every cycle. Each ALU Pipe can use one of the following paths at any time:

- The main unit which can execute up to two instructions in parallel, one of which must be a 32-bit float, the other either a 32-bit float or any integer. The supported native operations include multiply-and-add, multiply and addition. This unit can also perform packing and test operations on the results of these and a final move operation.

- A 16-bit float Sum-of-Products unit, which can execute two operations of the form `(a * b) OP (c * d)`, where `OP` is either add, subtract, min or max.

- A bitwise unit. The supported native operations include logic operations, logical shifts and arithmetic shifts. In one cycle it is possible to perform up to two shifts with a logical operation in between.

- A 32-bit complex floating-point unit that can execute a single operation. The supported native operations include reciprocal, logarithmic, exponent, (cardinal) sine operations and arc tangent. The complex unit may require more than a single instruction to execute these (for example, requiring range reduction before executing sines, etc.) but in general these are many times more efficient than a full "software" implementation.

The compiler implements other operations as routines on top of these, consuming multiple cycles.

A.1.   Example Device: 500MHz G6400

Table 3 shows the theoretical rates of throughput achievable for a 500MHz G6400 device, with 4 USCs, for various simple operations.

Table 3. Theoretical rates of throughput for a 500MHz G6400 device

Data type	Operation	Operations per instruction	Cycles per instruction	Theoretical throughput of 0.5GHz, 4xUSC G6400
`16-bit float`	Sum-Of-Products	6	1	(0.5 × 4 × 16 × 6) ÷ 1 = **192 GFLOPS**
`float`	Multiply-and-Add	4	1	(0.5 × 4 × 16 × 4) ÷ 1 = **128 GFLOPS**
`float`	Multiply	2	1	(0.5 × 4 × 16 × 2) ÷ 1 = **64 GFLOPS**
`float`	Add	2	1	(0.5 × 4 × 16 × 2) ÷ 1 = **64 GFLOPS**
`float`	Divide[A]	1	4	(0.5 × 4 × 16 × 1) ÷ 4 = **8 GFLOPS**
`float`	Divide[B]	1	2	(0.5 × 4 × 16 × 1) ÷ 2 = **16 GFLOPS**
`int`	Multiply-and-Add	2	1	(0.5 × 4 × 16 × 2) ÷ 1 = **64 GIOPS**

Data type	Operation	Operations per instruction	Cycles per instruction	Theoretical throughput of 0.5GHz, 4xUSC G6400
int	Multiply	1	1	$(0.5 \times 4 \times 16 \times 1) \div 1 =$ **32 GIOPS**
int	Add	1	1	$(0.5 \times 4 \times 16 \times 1) \div 1 =$ **32 GIOPS**
int	Divide	1	30	$(0.5 \times 4 \times 16 \times 1) \div 30 =$ **1.07 GIOPS**

A. By default, the compiler implements `float` division as 2 range reductions, followed by reciprocal and multiplication instructions, requiring 4 cycles.

B. For OpenGL GLSL shaders, OpenCL kernels with `-cl-finite-math-only` or `-cl-fast-relaxed-math`, or RenderScript kernels with `#pragma rs_fp_imprecise`, the compiler omits the range reduction, requiring only 2 cycles.

A.2. Measured Throughput

Actual measured peak throughput is likely to be less than the theoretical maximum, due to system overheads such as dynamic clock scaling, system bandwidth, rendering the user interface and driver/hardware setup costs. Actual kernels are also unlikely to generate only arithmetic operations, and may include control-flow instructions and memory accesses.

Appendix B.   Complex Operations

Table 4 lists the costs for standard single-precision, floating-point functions (for OpenCL, it assumes the `-cl-fast-relaxed-math` compilation flag).

A `low` cost represents less than 30 hardware instructions, and a `high` cost represents more than 100 instructions.

Table 4. Cost of executing standard single-precision, floating-point functions

Function	Cost
`x+y`	Low
`x*y`	Low
`x/y`	Low
`1/x`	Low
`rsqrt(x), 1/sqrt(x)`	Low
`sqrt(x)`	Low
`cbrt(x)`	Low
`hypot(x)`	Medium
`exp(x)`	Medium
`exp2(x)`	Low
`exp10(x)`	Medium
`expm1(x)`	Medium
`log(x)`	Low
`log2(x)`	Low
`log10(x)`	Low
`log1p(x)`	High
`sin(x)`	Low
`cos(x)`	Low
`tan(x)`	High
`sincos(x, cptr)`	Medium
`asin(x)`	Medium
`acos(x)`	Medium
`atan(x)`	Low
`atan2(y, x)`	Medium
`sinh(x)`	Medium
`cosh(x)`	Medium
`tanh(x)`	Medium
`asinh(x)`	High
`acosh(x)`	High
`atanh(x)`	High
`pow(x,y)`	High

Function	Cost
`fma(x, y, z)`	Low
`frexp(x,exp)`	Low
`ldexp(x, exp)`	Medium
`logb(x)`	Low
`ilogb(x)`	Low
`fmod(x, y)`	Medium
`remainder(x, y)`	High
`remquo(x, y, iptr)`	High
`modf(x, iptr)`	Low
`fdim(x, y)`	Low
`trunc(x)`	Low
`round(x)`	Low
`rint(x)`	Low
`ceil(x)`	Low
`floor(x)`	Low
`acospi(x)`	Medium
`asinpi(x)`	Medium
`atan2pi(x, y)`	Medium
`atanpi(x)`	Low
`copysign(x, y)`	Low
`cospi(x)`	Medium
`fabs(x)`	Low
`fmax(x, y)`	Low
`fmin(x, y)`	Low
`fract(x, ptr)`	Low
`hypot(x, y)`	Medium
`ldexp(x, y)`	Medium
`mad(x, y, z)`	Low
`maxmag(x, y)`	Low
`minmag(x, y)`	Low
`nan(x)`	Low
`nextafter(x, y)`	Low
`powr(x, y)`	High
`pown(x, y)`	High
`rootn(x, y)`	High
`signbit(x)`	Low
`sinpi(x)`	Medium

Function	Cost
`tanpi(x)`	Medium

Appendix C.   Optimization Strategy Cheat Sheet

Choose the correct device for your task

- Graphics Core: large datasets, simple flow, coherent data, large numbers of computations in general.

- CPU: complex flow control and branching, one to a few threads of execution, little parallelism, one-shot tasks.

Optimise your algorithm for convergence

- Avoid algorithms with complex flow control, or threads in a workgroup that are idle or exit early.

- Define and execute your kernel for maximum data coherency of nearby items.

- If having flow control, try to maximise the probability that all threads in a workgroup will take the same path.

If possible, define a workgroup size that is compile time known and a multiple of 32

- Occupancy, register pressure and many aggressive optimisations benefit from a workgroup size that is known at compile time and that is a multiple of 32.

- Hitting both of these conditions maximises benefits, but work singularly as well. Compile-time workgroup size enables several optimisations, while a workgroup that is a multiple of 32 maximises occupancy.

Use a large data set

- Provide enough work for the Graphics Core to schedule so that latency can be hidden.

- A few thousands to a few millions should be the target.

- Only consider less than a few thousands if the kernel is sufficiently long to warrant setting up.

Minimise bandwidth/maximise utilisation

- Leverage shared memory to minimise bandwidth use.

- Use a healthy amount of arithmetic instructions interspersed with memory reads/writes to hide latency.

- Schedule a lot of work simultaneously to assist the scheduler hide memory fetch latency.

- Consider using arithmetic instead of using lookup tables.

Maximise occupancy

- Try to use temporary storage (unified store) and shared memory (common store) sensibly so as not to hurt occupancy. Registers are not infinite! Up to ~40 32-bit registers per work-item should be sufficient.

- Minimise necessary temporaries.

- Use a healthy amount of arithmetic instructions interspersed with memory reads/writes to hide this latency. Consider using arithmetic versus lookups when bandwidth is limited.

Balance minimising bandwidth use with maximising occupancy

- Workgroup size, shared memory use and private memory use must all be tweaked together.

- Bandwidth should be the first target of optimisation in most applications.

- However, too much shared or private memory use may force the compiler to reduce utilisation.

Trade precision for speed wherever possible

- Consider using `half` instead of `float` (OpenCL), `mediump` instead of `highp` (OpenGL ES).

Access shared memory sequentially (on 2D accesses, in row-major order)

- Access sequential elements from kernel instances to maximise shared memory access speed and minimise bank conflicts.

- Do not use stride when accessing values. Strided access in most cases directly increases the number of conflicts. Maximum shared memory speed is achieved by accessing sequential elements.

Access raw global memory linearly and textures in square patterns

- Accessing memory linearly helps cache locality and burst transfers.

- Textures hardware is optimised to bring "nearby" pixels most efficiently.

Balance the length of the kernel: aim for a few tens to a few hundreds of hardware instructions

- If the kernel is not long enough, the cost of setting up might end up being a big part of overall execution time.

- Excessively long kernels (thousands of instructions or long loops) are generally efficient, but can starve other system components (notably Graphics).

Use shared objects with API Interoperability to avoid expensive redundant copies

- See what options and extensions are available to the combination of the APIs you are using (e.g. EGL Image).

- Using a shared image or buffer and avoiding a round-trip to the CPU will almost always prove a big gain.

Leverage burst transfers

- Help the compiler identify burst transfers by grouping reads and writes together making them easy to identify. Use explicit block-copy instructions whenever possible, such as the OpenCL `async_work_group_copy`.

Try to hide common-store latency

- If possible try to have a few unrelated calculations between accesses to shared memory and their corresponding barrier syncs.